

CSCE 790-002: Deep Reinforcement Learning and Search
Homework 3
Due: 10/28/2020

Installation

We will be using an updated `conda` environment.

To update, follow the same instructions from Homework 1 using the updated `spec-file.txt` file.

The entire GitHub repository should be downloaded again as changes were made to other files. You can download it with the green “Code” button and click “Download ZIP”.

1 Gradient Descent

For this exercise, you will implement forward propagation and backward propagation for a linear layer and a pointwise non-linear layer for a neural network. These two functions that you will implement will build your understanding of how neural networks work. The two layers you implement in this exercise can be composed in a variety of ways to build powerful function approximators.

There are N inputs of dimension D . The input matrix is $\mathbf{X} \in \mathbb{R}^{N \times D}$.
Where \mathbf{x}_n is the n^{th} input and x_{nd} is the n^{th} input at index d .

To check your implementations run:

```
python run_check_gradients.py --layer relu --num_inputs 5 --input_dim 10 --hidden_dim 100
python run_check_gradients.py --layer linear --num_inputs 5 --input_dim 10 --hidden_dim 100
```

The goal is to have the squared error between your implementation of the output computation step and the gradient computation be zero.

Vary `--num_inputs`, `--input_dim`, and `--hidden_dim` to verify your implementation.

1.1 Gradient of a Rectified Linear Unit (ReLU) Layer (10 pts)

Implement `relu_forward` and `relu_backward` in `assignments_code/assignment3.py`.

Given an input $\mathbf{X} \in \mathbb{R}^{N \times D}$, `relu_forward` returns the output $\mathbf{O} \in \mathbb{R}^{N \times D}$.
where $o_{nd} = \max(o_{nd}, 0)$.

Given a backpropagated gradient, $\delta \in \mathbb{R}^{N \times D}$, `relu_backward` backpropagates the gradient to the:

- inputs: $\frac{\partial E}{\partial x_{nd}} = \delta_{nd}(x_{nd} > 0)$.

1.2 Gradient of Linear Layer (40 pts)

Implement `linear_forward` and `linear_backward` in `assignments_code/assignment3.py`.
This linear layer takes as input a D dimensional vector and outputs a K dimensional vector.

There is a weight matrix $\mathbf{W} \in \mathbb{R}^{K \times D}$ and a bias vector $\mathbf{b} \in \mathbb{R}^K$.

`linear_forward` computes $\mathbf{O} \in \mathbb{R}^{N \times K} = \mathbf{X}\mathbf{W}^T + \mathbf{b}$
 $o_{nk} = \sum_{i=0}^D x_{ni}w_{ki} + b_k$

Given a backpropagated gradient, $\delta \in \mathbb{R}^{N \times K}$, `linear_backward` backpropagates the gradient to the:

- weights: $\nabla_{\mathbf{W}} E = \delta^T \mathbf{X}$, $\frac{\partial E}{\partial w_{kd}} = \sum_{n=0}^N \delta_{nk} x_{nd}$
- biases: $\frac{\partial E}{\partial b_k} = \sum_{n=0}^N \delta_{nk}$
- inputs: $\nabla_{\mathbf{X}} E = \delta \mathbf{W}$, $\frac{\partial E}{\partial x_{nd}} = \sum_{k=0}^K \delta_{nk} w_{kd}$

2 Deep Q-Network (DQN)

2.1 Implement a DQN (10 pts)

Implement `get_dqn` in `assignments_code/assignment3.py`.

The implementation should be in PyTorch. The DQN should take an input for dimension 100, followed by a hidden layer of size 75, followed by a hidden layer of size 50, and, finally, followed by an output layer of size 4. The hidden layers should have a ReLU activation function while the output layer should have a linear activation function. The output layer represents $\hat{q}(s, a, \mathbf{w})$ for all four actions.

2.2 Deep Q-learning (40 pts)

Implement `deep_q_learning_step` (shown in Algorithm 2) in `assignments_code/assignment3.py`.

Like in the previous homework about Q-learning, we will be implementing the Q-learning algorithm with a deep neural network. To see the answers to Homework 2, see Course Content - Homework - Homework2.

These commands have already been run:

```
torch.set_num_threads(1)
device: torch.device = torch.device("cpu")
dqn: nn.Module = get_dqn()
optimizer: Optimizer = optim.Adam(dqn.parameters(), lr=0.001)
```

Running the code:

```
python run_deep_q_learning.py --map maps/map1.txt --wait_greedy 0.1
```

Compare your results to the solution videos provided in Blackboard under Course Content - Homework - Homework 3. Keep in mind that this algorithm is stochastic, so results will not be exactly the same.

What to Turn In

Turn in your implementation of `assignments_code/assignment3.py`.

Algorithm 1 Deep Q-learning

```
1: procedure DEEP_Q-LEARNING( $\gamma, \epsilon, \alpha, B$ )  
2:   Initialize  $\hat{q}, \hat{q}_\tau$  ▷ DQN and target DQN  
3:    $D = []$  ▷ initialize replay buffer as empty list  
4:   for  $episode \in 1 \dots N$  do  
5:     Initialize  $s$   
6:     for  $t \in 1 \dots T$  do  
7:        $s, \hat{q}, D = \text{Deep\_Q\_learning\_Step}(s, \hat{q}, \hat{q}_\tau, \gamma, \epsilon, \alpha, B, D)$   
8:     end for  
9:     if Update then  
10:       $\hat{q}_\tau = \hat{q}$   
11:    end if  
12:  end for  
13:  return  $\hat{q}$  ▷ Approximation of  $q_*$   
14: end procedure
```

Algorithm 2 Deep Q-learning Step

```
1: procedure DEEP_Q-LEARNING-STEP( $s, \hat{q}, \hat{q}_\tau, \gamma, \epsilon, \alpha, B, D$ )  
2:   Sample an action  $a$  from  $\epsilon$ -greedy policy derived from  $Q$   
3:    $s', r = \text{env.sample\_transition}(s, a)$   
4:   Store transition  $(s, a, r, s')$  in  $D$   
5:   Sample a batch of data  $(\mathbf{X}, \mathbf{y})$  of size  $B$  from  $D$   
6:    $\mathbf{X} \in \mathbb{R}^{B \times d}$  where  $d$  is the dimension of the input to  $\hat{q}$   
7:    $\mathbf{y} \in \mathbb{R}^B$   
8:    $\mathbf{x}_i = \text{env.state\_to\_nnet\_input}(s_i)$   
9:    $\mathbf{x}'_i = \text{env.state\_to\_nnet\_input}(s'_i)$   
10:   $y_i = r_i + \gamma \max_{a'} \hat{q}_\tau(x'_i, a', \mathbf{w}^-)$   
11:  Calculate  $\nabla_{\mathbf{w}} E(\mathbf{w})$  with error  $E(\mathbf{w}) = \frac{1}{B} \sum_{i=1}^B (y_i - \hat{q}(x_i, a_i, \mathbf{w}))^2$   
12:  Update  $\mathbf{w}$  with learning rate  $\alpha$   
13:  return  $s', \hat{q}, D$   
14: end procedure
```
